

STAT 431 — Applied Bayesian Analysis — Course Notes

Bayesian Computation: Monte Carlo

Fall 2022

Numerical integration works well for low-dimensional problems (few parameters).

For high-dimensional problems, simulation is often better ...

Any method using randomized simulation (sampling) for approximation is called a **Monte Carlo** method.

Independent Sampling

Idea: Randomly generate samples from the posterior.

Then use the **empirical distribution** of the samples to estimate aspects of the actual posterior.

Let data be \mathbf{y} , and let the sample from the posterior for parameter θ be

$$\theta^{(1)}, \dots, \theta^{(S)}$$

A posterior sample may be used to approximate many things —

► a mean:

$$E(\theta \mid \mathbf{y}) \approx \frac{1}{S} \sum_{s=1}^S \theta^{(s)} = \text{sample mean of } \theta^{(s)}_s$$

► a variance:

$$\text{Var}(\theta \mid \mathbf{y}) \approx \text{sample variance of } \theta^{(s)}_s$$

► a (lower) quantile $Q_\theta(\tau)$:

[Draw quantile ...]

Round τS to the nearest integer: $[\tau S]$

Then use the $[\tau S]$ th order statistic from the sample.

(The order statistics are the $\theta^{(s)}$ s re-ordered from least to greatest.)

- ▶ a 95% equal-tailed credible interval: form an estimated version of

$$(Q_{\theta}(0.025), Q_{\theta}(0.975))$$

- ▶ probability of $H_0 : \theta \in \Theta_0$

$$\text{fraction of } \theta^{(s)} \text{s in } \Theta_0 = \frac{1}{S} \sum_{s=1}^S I(\theta^{(s)} \in \Theta_0)$$

where I represents the indicator function:

$$I(\theta \in \Theta_0) = \begin{cases} 1 & \theta \in \Theta_0 \\ 0 & \theta \notin \Theta_0 \end{cases}$$

- a mean of a function of θ :

$$\mathbb{E}(g(\theta) \mid \mathbf{y}) \approx \frac{1}{S} \sum_{s=1}^S g(\theta^{(s)})$$

This is valid because

$$g(\theta^{(1)}), \dots, g(\theta^{(S)})$$

is a random sample from the posterior of $g(\theta)$.

Since these approximations are subject to random sampling variability, we need an assessment of their accuracy.

For example, the approximation

$$\frac{1}{S} \sum_{s=1}^S g(\theta^{(s)}) \quad \text{of} \quad E(g(\theta) \mid \mathbf{y})$$

has an approximate *standard error* of

$$\frac{sd_g}{\sqrt{S}} \quad \text{where} \quad sd_g = \sqrt{\text{sample var. of } g(\theta^{(s)})}$$

(why?)

This is the **Monte Carlo error** for the mean approximation.

Example: Jevons's Coins Comparison

For $n_1 = 24$ coins minted before 1830:

$$\bar{y}_1 = 7.8730 \quad s_1 = \text{sample s.d.} = 0.05353$$

For $n_2 = 123$ newer coins (1860's):

$$\bar{y}_2 = 7.9725 \quad s_2 = \text{sample s.d.} = 0.01409$$

Assume independent samples from $\text{Normal}(\mu_1, \sigma_1^2)$ and $\text{Normal}(\mu_2, \sigma_2^2)$.

We want inference about $\mu_1 - \mu_2$.

Take “independent” “standard” (product-Jeffreys) priors:

$$\mu_1, \sigma_1^2, \mu_2, \sigma_2^2 \sim \frac{1}{\sigma_1^2} \cdot \frac{1}{\sigma_2^2} d\mu_1 d\sigma_1^2 d\mu_2 d\sigma_2^2$$

flat

Then it turns out that (later?) *as it doesn't comply with frequentist inference that's why this*

(μ_1, σ_1^2) and (μ_2, σ_2^2) are posterior-independent

with posterior distributions ($i = 1, 2$)

$$\underline{\mu_i \mid \sigma_i^2, \mathbf{y}} \sim \underline{\text{Normal}(\bar{y}_i, \sigma_i^2/n_i)}$$

$$\underline{\sigma_i^2 \mid \mathbf{y}} \sim \underline{\text{InvGamma}\left(\frac{n_i - 1}{2}, \frac{n_i - 1}{2} s_i^2\right)}$$

where \mathbf{y} represents both samples together.

strictly for positive posterior to be proper

As it turns out, under the posterior, μ_1 and μ_2 have independent (location-scale) t distributions. (Later?)

Nonetheless, $\mu_1 - \mu_2$ has no simple-form posterior density. (quantiles)
but 'mean, var' ✓

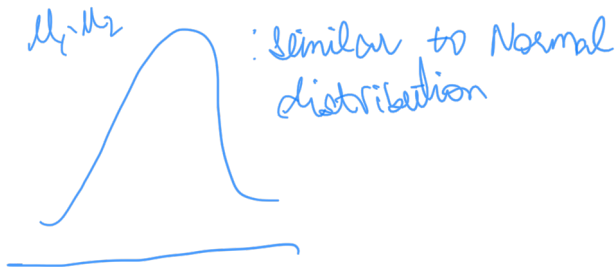

We can easily randomly sample from the posterior distribution of $\mu_1 - \mu_2$ using R ...

⇒ simulate μ_i

- 9nt
- DH
- ML
- 546
- MLD
- LC

R Example 3.3:

Comparing Normal Means: Independent Sampling



Barrett's Fischer Problem: Mean comparison of
2 Normal dist. with both assumed to have diff

The 95% Welch interval is an approximate frequentist confidence interval for $\mu_1 - \mu_2$, used here for comparison: variance

$$\bar{y}_1 - \bar{y}_2 \pm t_{0.025, df} \cdot \sqrt{\frac{s_1^2}{n_1} + \frac{s_2^2}{n_2}}$$

where

$$df = \frac{(s_1^2/n_1 + s_2^2/n_2)^2}{(s_1^2/n_1)^2/(n_1 - 1) + (s_2^2/n_2)^2/(n_2 - 1)}$$

Null:
 $H_0: \mu_1 \leq \mu_2$
frequentist F-test p-value
 $1 - pf\left(\frac{s_1^2}{s_2^2}, n_1 - 1, n_2 - 1\right)$
 $f \downarrow$ - distr.